

- 001.** In Iterator, nextElement() method of Enumeration has been changed to ----- **A**
 A next() B getNext()
 C returnNext() D name remains same
- 002.** TreeMap implements----- **D**
 A Dictionary B HashMap
 C AbstractMap D NavigableMap
- 003.** TreeSet internally uses which one to store elements? **C**
 A HashMap B LinkedHashMap
 C TreeMap D Tree
- 004.** HashSet internally uses **A**
 A HashMap B LinkedHashMap
 C TreeMap D Tree
- 005.** What guarantees type-safety in a collection? **A**
 A Generics B Abstract classes
 C Interfaces D Collection
- 006.** HashSet internally uses? **C**
 A Set B HashMap
 C List D Collection
- 007.** What implementation of Iterator can traverse a collection in both directions? **B**
 A Iterator B ListIterator
 C SetIterator D MapIterator
- 008.** Which does NOT implement the Collection interface? **B**
 A List B Map
 C Set D Array
- 009.** The Comparable interface contains method called----- **C**
 A to Compare B compare
 C compareTo D compareTo
- 010.** java.util.Collections is a: **A**
 A class B interface
 C package D object
- 011.** What is invalid about serialization in java? **B**
 A Serialized object can be persisted into database. B Serialized object can't be persisted into file.
 C Serialization is process of converting object into byte stream D Serialized object can be transferred over network.
- 012.** What method should we use for writing object in Serialization in java? **C**
 A objectWrite B objWrite
 C writeObject D write
- 013.** Can you customize serialization process when you have implemented Serializable interface **D**
 A By defining objectWrite() and objectRead() methods B It's not possible
 C By defining write() and read() methods D By defining writeObject() and readObject() methods
- 014.** Which of these provides a get(in index) method? **C**
 A Map B Set
 C List D Stack
- 015.** Which of these does NOT have an index based structure? **B**
 A Map B Set
 C List D Stack
- 016.** An attempt to add null key to a TreeSet will result in: **D**
 A Will compile B Compile time Exception
 C Error D Runtime - NullPointerException
- 017.** How can you sort given HashMap on basis of values ? **A**

- A Implement Comparator interface and override its compare method B It's not possible
- C Implement Comparator interface and override its compareTo method D Implement Comparator and Comparable interface
- 018.** Which interface must be implemented for sorting on basis many criteria? **A**
- A Comparator B Comparable
- C Serializable D Compare
- 019.** Which of the following are not a wrapper class? **B**
- A Double B String
- C Integer D Boolean
- 020.** Which of these is a super class of wrapper classes Long & Integer? **A**
- A Number B Int
- C Character D Long
- 021.** A -----lets you insert and remove elements at only one end **B**
- A queue B stack
- C list D set
- 022.** Which of these is a process of extracting/removing the state of an object from a stream? **D**
- A Serialization B Externalization
- C File Filtering D Deserialization
- 023.** All the wrapper classes (Integer, Boolean, Float, Short, Long, Double and Character) in Java are ----- **D**
- A immutable B private
- C serializable D final
- 024.** Which of these type parameters is used for a generic class to return and accept any type of object? **C**
- A K B N
- C T D V
- 025.** What will happen if we try to serialize List, Set and Map as member of class in java? **B**
- A NotSerializableException B They will get serialized
- C compilation error D SerializableNotPossibleException
- 026.** What will happen if we try to serialize primitive types in java? **A**
- A They will get serialized B compilation error
- C SerializableNotPossibleException D NotSerializableException
- 027.** Why are generics used? **C**
- A generics make code more fast B generics make code more optimized and readable
- C generics add stability to the code by making bugs detectable at compile time D generics add stability to the code by making bugs detectable at run time
- 028.** ----- involves the design and implementation of data structures and algorithms that work for multiple types **C**
- A top down programming B module programming
- C Generic programming D abstract programming
- 029.** A -----is a data type that keeps associations between keys and values **B**
- A set B map
- C tree D heap
- 030.** -----dont have duplicates. Adding a duplicate of an element that is already present is silently ignored. **A**
- A set B map
- C tree D heap
- 031.** Which of these is method of ObjectOutputStream interface used to write the object to input or output stream as required? **D**
- A write() B Write()

- C Byte D Double
- 048.** Which of the following statement is not correct about HashMap class? **D**
 A HashMap class extends AbstractMap B HashMap class is not synchronized and implements Map interface
 C HashMap uses a hashtable to store the map D HashMap maintain order of its element
- 049.** Which of the following method is used to return a Set that contains the entries in a map ? **C**
 A keyset() B getSet()
 C entrySet() D getAll()
- 050.** Which of these is a wrapper for simple data type char? **B**
 A Float B Character
 C String D Integer
- 051.** Which of the following methods is a method of wrapper Integer for obtaining hash code for the invoking object? **C**
 A int hash() B int hashCode()
 C int hashCode() D Integer hashCode()
- 052.** Which of the following is method of wrapper Integer for converting the value of an object into int? **B**
 A bytevalue() B int intValue()
 C Bytevalue() D Byte Bytevalue()
- 053.** What package is a part of the wrapper class which is imported by default into all Java programs? **A**
 A java.lang B java.awt
 C java.io D java.util
- 054.** Which of these is a process of writing the state of an object to a byte stream? **A**
 A Serialization B Externalization
 C File Filtering D Deserialization
- 055.** How will you implement dynamic arrays in Java? **D**
 A Set B Map
 C HashMap D List
- 056.** Which of the following is not a goal of HDFS? **C**
 A Fault detection and recovery B Handle huge dataset
 C Prevent deletion of data D Provide high network bandwidth for data movement
- 057.** In which mode each daemon runs on a single node but there is separate java process for each daemon **D**
 A local B standalone
 C fully distributed D pseudo distributed
- 058.** HDFS refers to **D**
 A Hadoop Distributed Flat file System B Hadoop Distributive File System
 C Hadoop Data File System D Hadoop Distributed File System
- 059.** Which of the following interface guarantees that no duplicates elements are stored and all elements are accessed in natural order ? **B**
 A java.util.List B java.util.Set
 C java.util.Deque D java.util.Map
- 060.** What does get(int index) method define by List interface do ? **A**
 A returns an object stored at the specified index B stores an object at the specified index
 C returns a list containing elements between specified index and end in the collection D returns index
- 061.** Which interface provides the capability to store objects using a key- value pair ? **A**
 A java.util.Map B java.util.List

- C java.util.Set D java.util.Collection
- 062.** Which interface is used to traverse a list in both forward and backward direction? **B**
 A Iterator interface B ListIterator interface
 C Enumeration interface D DoubleList Iterator
- 063.** Which of the following statement is true about TreeSet class ? **C**
 A TreeSet extends AbstractList class and implements the List interface B Uses a LinkedList structure to store elements
 C TreeSet stores elements sorted ascending order D Access and retrieval times are quite slow
- 064.** The logical division of data is known as -----while a physical division of data is known as ----- **D**
 A record, block B split, record
 C HDFS block, split D Split, HDFS Block
- 065.** The process by which the system performs the sort and transfers the map outputs to the reducer as inputs is known as the **B**
 A reduce B shuffle
 C sort D combine
- 066.** The HDFS command to create the copy of a file from a local system is which of the following? **A**
 A copyFromLocal B copyfromlocal
 C CopyFromLocal D copyLocal
- 067.** A GFS cluster consists of a -----master and -----chunk-servers and is accessed by multiple clients **B**
 A one , one B single, multiple
 C multiple, multiple D multiple, single
- 068.** Rack awareness is the way in which the -----determines on how to place blocks based on the rack definitions. **D**
 A job tracker B secondary name node
 C data node D namenode
- 069.** When using HDFS, what occurs when a file is deleted from the command line? **C**
 A It is permanently deleted if trash is enabled. B It is placed into a trash directory common to all users for that cluster.
 C It is permanently deleted and the file attributes are recorded in a log file. D It is moved into the trash directory of the user who deleted it if trash is enabled
- 070.** The information mapping data blocks with their corresponding files is stored in **A**
 A namenode B datanode
 C job tracker D task tracker
- 071.** GFS refers to **B**
 A Google Fit file System B Google File System
 C Google Flat file System D Google Fixed file System
- 072.** The fully distributed mode of installation (without virtualization) needs a minimum of ----- physical machines **C**
 A 1 B 2
 C 3 D 4
- 073.** When a jobTracker schedules a task is first looks for **A**
 A A node with empty slot in the same rack as datanode B Any node on the same rack as the datanode
 C Any node on the rack adjacent to rack of the datanode D Just any node in the cluster
- 074.** Which of the following is/are true about combiners? **D**
 A Combiners can be used for mapper only job B Combiners can be used for any Map Reduce operation
 C Mappers can be used as a combiner D Combiners are primarily aimed to

- class improve Map Reduce performance
- 075.** Which of the following is a common hadoop maintenance issue? **B**
- A Lack of tools B Lack of configuration management
C Lack of web interface D lack of memory
- 076.** Which of following statement(s) are correct? **C**
- A Master and slaves files are optional in Hadoop 2.x B Master file has list of all name nodes
C Core-site has hdfs and MapReduce related common properties D hdfs-site file is now deprecated in Hadoop 2.x
- 077.** What is the default HDFS replication factor? **C**
- A 4 B 1
C 3 D 2
- 078.** The current limiting factor to the size of a hadoop cluster is **C**
- A Excess heat generated in data center B Upper limit of the network bandwidth
C Upper limit of the RAM in namenode D 4000 data nodes
- 079.** When a job tracker schedules a task , it first looks for **C**
- A any node on the rack adjacent to rack of datanode B any node on the same rack as datanode
C a node with empty slot in the same rack as datanode D any node in the cluster
- 080.** Master node in GFS, periodically communicates with each chunkserver in ----- messages **B**
- A RPC B HeartBeat
C IPC D text
- 081.** By default, the GFS makes ---- replicas per chunk **C**
- A 2 B 4
C 3 D 1
- 082.** Which of the following can be used to control the number of part files in a map reduce program output directory? **B**
- A Number of Mappers B Number of Reducers
C Counter D Partitioner
- 083.** Which of the following is NOT a type of metadata in NameNode? **C**
- A List of files B Block locations of files
C No. of file records D File access control information
- 084.** What does Velocity in Big Data mean? **D**
- A Speed of input data generation B Speed of individual machine processors
C Speed of ONLY storing data D Speed of storing and processing data
- 085.** _____ has the worlds largest Hadoop cluster. **C**
- A Apple B Datamatics
C Facebook D Microsoft
- 086.** Which of the following are true for Hadoop Pseudo Distributed Mode? **C**
- A It runs on multiple machines B Runs on multiple machines without any daemons
C Runs on Single Machine with all daemons D Runs on Single Machine without all daemons
- 087.** Which of the following command runs the HDFS secondary namenode? **B**
- A secondary namenode B secondarynamenode
C secondary_namenode D secondary-namenode
- 088.** A file in HDFS that is smaller than a single block size **C**
- A Cannot be stored in HDFS. B Occupies the full block 's size.
C Occupies only the size it needs and not the full block D Can span over multiple blocks.
- 089.** The command `hadoop fs -test -z URI` gives the result 0 if **D**

- A if the path is a directory
C if the path is not empty
- B if the path is a file
D if the file is zero length
- 090.** The hdfs command to create the copy of a file from a local system is **D**
A CopyFromLocal
B copyfromlocal
C CopyLocal
D copyFromLocal
- 091.** In order to read any file in HDFS, instance of _____ is required. **A**
A filesystem
B datastream
C outstream
D inputstream
- 092.** Which one of the following stores data? **B**
A Name node
B Data node
C Master node
D task tracker
- 093.** Which hdfs command is used to check for various inconsistencies? **B**
A fsk
B fsck
C fetchdt
D sck
- 094.** HDFS block size is larger as compared to the size of the disk blocks so that **A**
A A single file larger than the disk size can be stored across many disks in the cluster.
B Transfer of a large files made of multiple disk blocks is not possible.
C The seek time is maximum
D Only HDFS files can be stored in the disk used
- 095.** The split size is normally the size of an _____ block, which is appropriate for most applications. **D**
A Generic
B Task
C Library
D HDFS
- 096.** For YARN, the _____ Manager UI provides host and port information. **C**
A Data Node
B NameNode
C Resource
D Replication
- 097.** HDFS supports the _____ command to fetch Delegation Token and store it in a file on the local system. **B**
A fetdt
B fetchdt
C fsk
D rec
- 098.** In _____ mode, the NameNode will interactively prompt you at the command line about possible courses of action you can take to recover your data. **C**
A full
B partial
C recovery
D commit
- 099.** HDFS is implemented in _____ programming language. **B**
A C++
B Java
C Scala
D SQL
- 100.** When the namenode finds that some blocks are over replicated, it **C**
A Stops the replication job in the entire hdfs file system.
B It slows down the replication process for those blocks
C It deletes the extra blocks.
D It leaves the extra blocks as it is.
- 101.** The following is not permitted on HDFS files **D**
A Deleting
B Renaming
C Moving
D Executing
- 102.** The client reading the data from HDFS filesystem in Hadoop does which of the following? **A**
A Gets only the block locations from the namenode
B Gets the data from the namenode
C Gets both the data and block location from the namenode
D Gets the block location from the datanode
- 103.** Which configuration file is used to control the HDFS replication factor? (S) **C**
A mapred-site.xml
B core-site.xml
C hdfs-site.xml
D yarn-site.xml

- 104.** A _____ serves as the master and there is only one NameNode per cluster. **B**
 A Data Node B NameNode
 C Data block D Replication
- 105.** A _____ node acts as the Slave and is responsible for executing a Task assigned to it by the JobTracker. **C**
 A MapReduce B Mapper
 C TaskTracker D JobTracker
- 106.** HDFS provides a command line interface called _____ used to interact with HDFS. **B**
 A HDFS Shell B FS Shell
 C DFS Shell D vi shell
- 107.** Which of the following is used for the MapReduce job Tracker node ? **C**
 A mradmin B tasktracker
 C jobtracker D job-tracker
- 108.** _____ is the slave/worker node and holds the user data in the form of Data Blocks. **A**
 A DataNode B NameNode
 C Data block D Replication
- 109.** Which command is used to show all the Hadoop daemons that are running on the machine **B**
 A distcp B jps
 C dir D fsck
- 110.** All the files in a directory in HDFS can be merged together using which of the following? **A**
 A Put merge B Get merge
 C Remerge D Merge all
- 111.** What does commodity Hardware in Hadoop world mean? **D**
 A Very cheap hardware B Industry standard hardware
 C Discarded hardware D Low specifications Industry grade hardware
- 112.** Which of the following is not a valid Hadoop config file? **C**
 A core-default.xml B hdfs-default.xml
 C hadoop-default.xml D mapred-default.xml
- 113.** Point out the wrong statement : **D**
 A Replication Factor can be configured at a cluster level (Default is set to 3) and also at a file level B Block Report from each DataNode contains a list of all the blocks that are stored on that DataNode
 C User data is stored on the local file system of DataNodes D DataNode is aware of the files to which the blocks stored on it belong
- 114.** Which of the following platforms does Apache Hadoop run on ? **C**
 A Dos like B Unix like
 C cross platform D debian
- 115.** In which mode all daemons execute in separate nodes **C**
 A local mode B pseudo distributed mode
 C fully distributed mode D stand alone mode
- 116.** All of the following accurately describe Hadoop, EXCEPT **B**
 A Open source B Real time
 C Java-based D Distributed computing approach
- 117.** Apache Hadoop achieves reliability by replicating the data across multiple hosts, and hence does not require _____ storage on hosts. **B**
 A Standard RAID levels B RAID
 C ZFS D Operating system
- 118.** When You are developing a combiner that takes as input Text keys, IntWritable values, and emits Text keys, IntWritable values. Which interface should your class implement? **B**
 A Combiner B Reducer

- 119.** The total number of partitioners is equal to -----

A Number of reducers B number of nodes
C number of combiners D number of mappers

120. Hadoop is a framework that allows the distributed processing of:

A Small Data Sets B Semi-Large Data Sets
C Large Data Sets D Large and Small Data sets

121. Which of the following is a valid flow in Hadoop ?

A Input -> Reducer -> Mapper -> Combiner -> -> Output B Input -> Mapper -> Reducer -> Combiner -> Output
C Input -> Mapper -> Combiner -> Reducer -> Output D Input -> Reducer -> Combiner -> Mapper -> Output

122. Which of the following method used to set the output directory?

A FileOutputFormat.setOutputgetpath() B OutputFormat.setOutputpath()
C FileOutputFormat.setOutputpath() D OutputFormat.setOutputgetpath()

123. Which of the following is/are true about combiners?

A Combiners can be used for mapper only job B Combiners can be used for any Map Reduce operation
C Mappers can be used as a combiner class D Combiners are primarily aimed to improve Map Reduce performance

124. Hadoop Framework is written in

A Java B Python
C C++ D Scala

125. Which of the following Hadoop config files is used to define the heap size?

A hdfs-site.xml B core-site.xml
C hadoop-env.sh D mapred-site.xml

126. Which of the following property gets configured on mapred-site.xml ?

A replication factor B java environment variables
C directory names to store hdfs files D host and port where MapReduce job runs

127. What is the default HDFS replication factor?

A 4 B 1
C 3 D 2

128. _____ can best be described as a programming model used to develop Hadoop-based applications that can process massive amounts of data.

A MapReduce B Mahout
C Oozie D Unix

129. _____ maps input key/value pairs to a set of intermediate key/value pairs.

A shuffle B Reducer
C Both Mapper and Reducer D Mapper

130. What was Hadoop named after?

A Creator Doug Cuttings favorite circus act B Cuttings high school rock band
C The toy elephant of Cuttings son D A sound Cuttings laptop made during Hadoop development

131. MapReduce was devised by ...

A Apple B Google
C Microsoft D Samsung

132. Number of mappers is decided by the

A Mappers specified by the programmer B Available Mapper slots
C Available heap memory D Input Splits

133. What is the default HDFS block size?

A 32 MB B 64 KB
C 128 KB D 64 MB

- 134.** What decides number of Mappers for a MapReduce job? **C**
 A File Location B mapred.map.tasks parameter
 C Input file size D Input Splits
- 135.** Above the file systems comes the _____ engine, which consists of one Job Tracker, to which client applications submit MapReduce jobs. **A**
 A MapReduce B Google
 C Functional programming D Facebook
- 136.** _____ part of the MapReduce is responsible for processing one or more chunks of data and producing the output results. **A**
 A Maptask B Mapper
 C Task execution D Shuffle
- 137.** _____ refers to incremental costs with no major impact on solution design, performance and complexity. **C**
 A Scale-out B Scale-down
 C Scale-up D Scale-in
- 138.** _____ is a generalization of the facility provided by the MapReduce framework to collect data output by the Mapper or the Reducer **B**
 A Partitioner B OutputCollector
 C Reporter D Collector
- 139.** Output of the mapper is first written on the local disk for sorting and _____ process. **A**
 A shuffling B secondary sorting
 C forking D reducing
- 140.** Mapper and Reducer implementations can use the _____ to report progress or just indicate that they are alive. **C**
 A Partitioner B OutputCollector
 C Reporter D Combiner
- 141.** _____ is the primary interface for a user to describe a MapReduce job to the Hadoop framework for execution. **B**
 A Map Parameters B JobConf
 C MemoryConf D Conf
- 142.** Which of the following phases occur simultaneously? **D**
 A Map and Reduce B Reduce and Sort
 C Shuffle and Map D Shuffle and Sort
- 143.** Mapper implementations are passed the JobConf for the job via the _____ method **B**
 A JobConfigure.configure B JobConfigurable.configure
 C JobConfigurable.configureable D Job.conf
- 144.** _____ is general-purpose computing model and runtime system for distributed data analytics. **A**
 A Mapreduce B Drill
 C Oozie D SQL
- 145.** _____ controls the partitioning of the keys of the intermediate map-outputs. **B**
 A Collector B Partitioner
 C InputFormat D Combiner
- 146.** The JobTracker pushes work out to available _____ nodes in the cluster, striving to keep the work as close to the data as possible. **B**
 A DataNodes B TaskTracker
 C ActionNodes D reduce-tracker
- 147.** The daemons associated with the MapReduce phase are _____ and task-trackers. **A**
 A job-tracker B map-tracker
 C reduce-tracker D reduce-tracker
- 148.** A Map reduce job can be written in: **D**
 A Java B Ruby
 C Python D Any Language which can read from

input stream

- 149.** _____ is a programming model designed for processing large volumes of data in parallel by dividing the work into a set of independent tasks. **B**
A Hive B MapReduce
C Pig D Lucene
- 150.** The output of the reduce task is typically written to the File System via ----- **A**
A OutputCollector B InputCollector
C OutputCollect D InputCollect
- 151.** _____ function is responsible for consolidating the results produced by each of the Map() functions/tasks. **A**
A Reduce B Map
C Reducer D sum
- 152.** Which of the following parameter is to collect keys and combined values? **D**
A key B values
C reporter D output
- 153.** Input to the _____ is the sorted output of the mappers. **A**
A Reducer B Mapper
C Shuffle D conf
- 154.** Which one of the following statements is false regarding a MapReduce job? **C**
A You can specify the number of Reducers using the setNumReduceTasks method of the JobConf object. B During the shuffling, the records are sorted by key and the values are combined into a collection.
C You can specify the number of Mappers using the setNumMapTasks method of the JobConf object. D All records that share the same key are sent to the same Reducer.
- 155.** Which MapReduce stage serves as a barrier, where all previous stages must be completed before it may proceed? **A**
A Combine B shuffle
C Reduce D Write
- 156.** Which MapReduce phase is theoretically able to utilize features of the underlying file system in order to optimize parallel execution? **A**
A Split B Combine
C Map D sort
- 157.** The new Hadoop API is in the -----package **B**
A org.apache.hadoop.mapred B org.apache.hadoop.mapreduce
C org.apache.hadoop.maprer D org.apache.hadoop.map
- 158.** In GFS, Chunkservers store the chunks on local disks as ----- **B**
A Dos files B Linux files
C Text file D Binary file
- 159.** What are the steps in the lifecycle of mapper? **B**
A JobTracker calls the map method of the mapper B Jobtracker calls the task tracker and task tracker calls the map method
C Task tracker calls the job tracker, which in turn calls the map method D job tracker and map are parallelly called
- 160.** Which of the following class is responsible for converting inputs to key-value Pairs of Map Reduce? **C**
A FileInputFormat B InputSplit
C RecordReader D Mapper
- 161.** When You are developing a combiner that takes as input Text keys, IntWritable values, and emits Text keys, IntWritable values. Which interface should your class implement? **B**
A Combiner <Text, IntWritable, Text, IntWritable> B Reducer <Text, IntWritable, Text, IntWritable>

- C Combiner <Text,Text, IntWritable, IntWritable> D Combiner <Text, Text, IntWritable, IntWritable>
- 162.** How are keys and values presented and passed to the reducers during a standard sort and shuffle phase of MapReduce? **A**
- A Keys are presented to reducer in sorted order; values for a given key are not sorted B Keys are presented to reducer in sorted order; values for a given key are sorted in ascending order
- C Keys are presented to a reducer in random order; values for a given key are not sorted D Keys are presented to a reducer in random order; values for a given key are sorted in ascending order
- 163.** How many instances of Job tracker can run on Hadoop cluster ? **A**
- A 1 B 2
- C 3 D 4
- 164.** ----- is also known as semi reducer. **B**
- A mapper B combiner
- C shuffler D sorter
- 165.** The number of partitioners is equal to the number of ----- **D**
- A nodes B mapper
- C combiner D reducers
- 166.** Which among the following controls the partitioning of the keys of the intermediate map-outputs? **C**
- A Reader B Combiner
- C Partitioner D Writer
- 167.** Job control is performed through the -----class in the new API, **B**
- A JobConf B Job
- C JobClient D Config
- 168.** Which of the following is not the Dameon process that runs on a hadoop cluster ? **C**
- A JobTracker B DataNode
- C TaskTracker D TaskNode
- 169.** In the old API both map and reduce outputs are named as **A**
- A part-nnnnn B part-nnnn
- C part-nnnnnn D part-nnn
- 170.** Calling the combiner function zero, one, or many times should produce -----output from the reducer **B**
- A zero B the same
- C null D key
- 171.** MapReduce programming model is ----- **B**
- A Platform Dependent but not language-specific B Neither platform- nor language-specific
- C Platform independent but language-specific D Platform Dependent and language-specific
- 172.** In a MapReduce job, the reducer receives all values associated with same key. Which statement best describes the ordering of these values? **B**
- A The values are in sorted order. B The values are arbitrarily ordered, and the ordering may vary from run to run of the same MapReduce job.
- C The values are arbitrary ordered, but multiple runs of the same MapReduce job will always have the same ordering. D Since the values come from mapper outputs, the reducers will receive contiguous sections of sorted values
- 173.** How can you disable the reduce step? **C**
- A The Hadoop administrator has to set the number of the reducer slot to zero on all slave nodes. This will disable B It is imposible to disable the reduce step since it is critical part of the Mep-Reduce abstraction.

the reduce step.

C A developer can always set the number of the reducers to zero. That will completely disable the reduce step.

D While you cannot completely disable reducers you can set output to one. There needs to be at least one reduce step in Map-Reduce abstraction

174. What are map files and why are they important?

D

A Map files are stored on the namenode B and capture the metadata for all blocks on a particular rack. This is how Hadoop is "rack aware"

Map files are the files that show how the data is distributed in the Hadoop cluster.

C Map files are generated by Map-Reduce after the reduce step. They show the task distribution during job execution

D Map files are sorted sequence files that also have an index. The index allows fast data look up.

175. What are supported programming languages for Map Reduce?

A

A The most common programming language is Java, but scripting languages are also supported via Hadoop streaming.

B Any programming language that can comply with Map Reduce concept can be supported.

C Only Java supported since Hadoop was written in Java.

D Currently Map Reduce supports Java, C, C++ and COBOL.

176. The -----command is used to check if the Hadoop daemons are running properly or not

B

A job

B jps

C dir

D ls

177. Your client application submits a MapReduce job to your Hadoop cluster. Identify the Hadoop daemon on which the Hadoop framework will look for an available slot schedule a MapReduce operation.

D

A TaskTracker

B NameNode

C DataNode

D JobTracker

178. Which of the following is the default partitioner for partitioning key space?

B

A Part

B Hash partitioner

C Partitioner

D Hpartitioner

179. The number of maps is usually driven by the total size of

A

A inputs

B tasks

C jobs

D outputs

180. What happens if number of reducers are set to 0?

D

A nothing happens

B reducer output will be final output

C Reduce only job takes place

D Map only job takes place